

Tanisha Jawale -UCE2022525

Srushti Kage-UCE2022531

Problem Statement: Prediction of CO2 Emissions from Vehicles

INTRODUCTIONS:

As concerns about environmental sustainability continue to grow, there is a pressing need to mitigate greenhouse gas emissions, particularly from transportation sources. Among these emissions, carbon dioxide (CO2) from vehicles contributes significantly to air pollution and climate change. Therefore, accurately predicting CO2 emissions from vehicles is crucial for policymakers, manufacturers, and consumers to implement effective measures to reduce environmental impact.

OBJECTIVE:

The goal of this project is to develop a predictive model that can accurately estimate CO2 emissions from vehicles based on various features such as vehicle type, engine specifications, fuel type, and mileage. By leveraging machine learning techniques, we aim to create models that not only predict CO2 emissions with high accuracy but also provide insights into the factors influencing emissions.

APPROACH:

Data Collection and Exploration:

- Gather a comprehensive dataset containing information on vehicle attributes and corresponding CO2 emissions.
- Explore the dataset to understand its structure, identify potential challenges (e.g., missing data, outliers), and gain insights into the relationships between features and CO2 emissions.

Data Preprocessing:

- Handle missing values, outliers, and any inconsistencies in the data.
- Encode categorical variables and normalize numerical features to prepare the data for modeling.

Feature Selection/Reduction:

- Employ feature selection techniques to identify the most relevant features for predicting CO2 emissions.
- Utilize dimensionality reduction methods if necessary to reduce the complexity of the dataset while preserving important information.

Model Selection and Training:

- Select a diverse set of regression models including Multiple Linear Regression, Lasso Regression, Decision Tree Regression, Ridge Regression, and Random Forest Regression.
- Train each model using the preprocessed data and evaluate their performance using appropriate metrics.

Model Evaluation and Interpretation:

- Assess the predictive performance of each model using metrics such as Mean Squared Error, R-squared, and Mean Absolute Error.
- Interpret the coefficients or feature importances of the models to understand the impact of different factors on CO2 emissions.
-

Prediction and Visualization:

- Use the trained models to make predictions on new or unseen data.
- Visualize the predicted CO2 emissions alongside actual values to assess the accuracy of the models and identify any patterns or trends.

CODE:

```
# Importing data & Checking for NULL Values

import pandas as pd
import numpy as np
```

```

dF=pd.read_csv('/content/CO2 Emissions_Canada.csv')
missing_vals=dF.isnull().sum()
print(missing_vals)
#checking for distinct company names

dF['Make'].unique()
#checking for distinct model names

dF['Model'].unique()
import numpy as np

# Putting different transmission sub-catalogies into their respective
catalogies

dF['Transmission']=np.where(dF['Transmission'].isin(['A4','A5','A6','A7','A8','A9','A10']),'Automatic',dF['Transmission'])
dF['Transmission']=np.where(dF['Transmission'].isin(['AV','AV7','AV8','AV6','AV10']),'Autonomous',dF['Transmission'])
dF['Transmission']=np.where(dF['Transmission'].isin(['AS4','AS5','AS6','AS8','AS7','AS9','AS10']),'Automatic Selection',dF['Transmission'])
dF['Transmission']=np.where(dF['Transmission'].isin(['M6','M5','M7']),'Manual',dF['Transmission'])
dF['Transmission']=np.where(dF['Transmission'].isin(['AM5','AM6','AM7','AM8','AM9','AM10']),'Automated Manual',dF['Transmission'])
print(dF['Transmission'].unique())
print()
print(dF['Transmission'].value_counts())

#finding correlation
print(dF.dtypes)
print()

# Check for non-numeric values in 'CO2 Emissions(g/km)' (binary values)
non_numeric_co2 = dF[~dF['CO2 Emissions(g/km)'].apply(lambda x:
isinstance(x, (int, float)))]
print(non_numeric_co2)
print()

# Check for non-numeric values in all columns

```

```

non_numeric_values = dF.applymap(lambda x: isinstance(x, str))
print(non_numeric_values.any())
print()

# Exclude non-numeric columns
numeric_columns = dF.select_dtypes(include=['float64', 'int64'])

# Calculate correlation of numeric columns w.r.t CO2
correlation = numeric_columns.corr()['CO2 Emissions(g/km)'].sort_values()
print(correlation)

import matplotlib.pyplot as plt
import seaborn as sns

# Calculate correlation
correlation = numeric_columns.corr()['CO2
Emissions(g/km)'].sort_values(ascending=False)

# Convert correlation Series to DataFrame with one column
correlation_df = correlation.to_frame()

plt.rcParams['figure.figsize'] = (10, 8)
sns.heatmap(correlation_df, cmap='coolwarm', linewidth=0.5, annot=True)

plt.title('Correlation')
plt.show()
# DATA VISUALIZATION
# Frequency distribution of cars of diff. companies

import matplotlib.pyplot as plt

plt.figure(figsize=(20,5))

dF.groupby('Make')['Make'].count().sort_values(ascending=False).plot(kind=
'bar',color='red')

plt.title('Frequency distribution of cars of different companies',
fontsize=25)
plt.xlabel('Company', fontsize=20)
plt.ylabel('Frequency', fontsize=20)

```

```
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
# Frequency distribution of models

import matplotlib.pyplot as plt

plt.figure(figsize=(20,5))

dF.groupby('Model')['Model'].count().sort_values(ascending=False)[:25].plot(kind='bar')

plt.title('Distribution of models', fontsize=25)
plt.xlabel('Models', fontsize=20)
plt.ylabel('Frequency', fontsize=20)
plt.xticks(rotation=90)
plt.tight_layout()
plt.show()

# Frequency distribution of vehicle class

import matplotlib.pyplot as plt

plt.figure(figsize=(20,5))

dF.groupby('Vehicle Class')['Vehicle
Class'].count().sort_values(ascending=False).plot(kind='bar',
color='green')

plt.title('Vehicle class distribution', fontsize=25)
plt.xlabel('Vehicle Class', fontsize=20)
plt.ylabel('Frequency', fontsize=20)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Frequency distribution of transmission

import matplotlib.pyplot as plt
import pandas as pd

dF = pd.read_csv('/content/CO2 Emissions_Canada.csv')
```

```

plt.figure(figsize=(20,5))

dF.groupby('Transmission')['Transmission'].count().sort_values(ascending=False).plot(kind='bar', color='magenta')

plt.title('Distribution of transmission', fontsize=25)
plt.xlabel('Transmission type', fontsize=20)
plt.ylabel('Frequency', fontsize=20)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

#checking for distinct model names

dF['Model'].unique()
#checking for distinct company names

dF['Make'].unique()
import numpy as np

# Putting different transmission sub-catagories into their respective
catagories

dF['Transmission']=np.where(dF['Transmission'].isin(['A4','A5','A6','A7','A8','A9','A10']),'Automatic',dF['Transmission'])
dF['Transmission']=np.where(dF['Transmission'].isin(['AV','AV7','AV8','AV6','AV10']),'Autonomous',dF['Transmission'])
dF['Transmission']=np.where(dF['Transmission'].isin(['AS4','AS5','AS6','AS8','AS7','AS9','AS10']),'Automatic Selection',dF['Transmission'])
dF['Transmission']=np.where(dF['Transmission'].isin(['M6','M5','M7']),'Manual',dF['Transmission'])
dF['Transmission']=np.where(dF['Transmission'].isin(['AM5','AM6','AM7','AM8','AM9','AM10']),'Automated Manual',dF['Transmission'])

print(dF['Transmission'].unique())
print()
print(dF['Transmission'].value_counts())
# Renaming fuel types for better understanding

dF['Fuel Type']= np.where(dF['Fuel Type']=='X', 'Regular gasoline', dF['Fuel Type'])

```

```

dF['Fuel Type']= np.where(dF['Fuel Type']=='Z', 'Premium gasoline',dF['Fuel Type'])

dF['Fuel Type']= np.where(dF['Fuel Type']=='E', 'Ethanol',dF['Fuel Type'])

dF['Fuel Type']= np.where(dF['Fuel Type']=='D', 'Diesel',dF['Fuel Type'])

dF['Fuel Type']= np.where(dF['Fuel Type']=='N', 'Natural gas',dF['Fuel Type'])

print(dF['Fuel Type'].unique())
print(dF['Vehicle Class'].unique())

#Counting no. of values of diff. fule types

dF['Fuel Type'].value_counts()

# Checking for rows where the fule type is natural gas & store the row index in ind

dF_N = dF[dF['Fuel Type']== 'Natural gas']

ind = dF_N.index

print(dF_N)
# Dropping the rows where the fule type is natural gas

for i in ind:
    dF.drop(i, axis=0, inplace=True)

import pandas as pd

d_v =pd.get_dummies(dF['Fuel Type'], prefix='Fuel', drop_first=True)
dv = pd.get_dummies(dF["Transmission"], drop_first=True)
d_v.head()

#Concatinating all 3 dataframes

df = [dF, d_v,dv]

data = pd.concat(df, axis=1)
data.head()

#Dropping the transmission & fule type columns from data(cancatinated dF)

```

```

df = [dF, d_v,dv]
data = pd.concat(df, axis=1)
data.drop(['Fuel Type'], inplace=True, axis=1)
data.drop(['Transmission'], inplace=True, axis=1)

#It replaces each value in the Make/Model/Vehicle_Class column with its
corresponding frequency stored in dict
# converts the categorical values into numerical representations

df_freq = data['Make'].value_counts().to_dict()
mod_freq = data['Model'].value_counts().to_dict()
veh_freq = data['Vehicle Class'].value_counts().to_dict()
data['Make'] = data['Make'].map(df_freq)
data['Model'] = data['Model'].map(mod_freq)
data['Vehicle Class'] = data['Vehicle Class'].map(veh_freq)
#Divide the dF into target & features for prediction purpose

X = data.drop('CO2 Emissions(g/km)', axis=1) # X contains all cols.
except CO2 (features)
y = data['CO2 Emissions(g/km)'] # Y conatains the CO2 col. (target)
# MULTIPLE REGRESSION

```

```

import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
import matplotlib.pyplot as plt

scaler =StandardScaler()

X = data.drop('CO2 Emissions(g/km)', axis=1)
y = data['CO2 Emissions(g/km)']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

regressor = LinearRegression()

```

```

regressor.fit(X_train, y_train)
my_dict={}
y_pred = regressor.predict(X_test)
y_pred_train = regressor.predict(X_train)

r2 = r2_score(y_test, y_pred)

my_dict['ML']=r2
# Plotting of train set
plt.title("CO2 Emissions (Training Set)")
plt.xlabel("Actual CO2 Emissions (g/km)")
plt.ylabel("Predicted CO2 Emissions (g/km)")
plt.scatter(y_train, y_pred_train, color="lightcoral", label="Training Data")
plt.plot(y_train, y_train, color="firebrick", label="Perfect Prediction")
plt.show()

# Plotting of test set
plt.title("CO2 Emissions (Testing Set)")
plt.xlabel("Actual CO2 Emissions (g/km)")
plt.ylabel("Predicted CO2 Emissions (g/km)")
plt.scatter(y_test, y_pred, color="lightcoral", label="Testing Data")
plt.plot(y_test, y_test, color="firebrick", label="Perfect Prediction")
plt.show()
# LASSO REGRESSION

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Lasso
from sklearn.metrics import r2_score
import matplotlib.pyplot as plt

X = data.drop('CO2 Emissions(g/km)', axis=1)
y = data['CO2 Emissions(g/km)']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

```

```

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

lasso = Lasso(alpha=0.1)
lasso.fit(X_train, y_train)

y_pred = lasso.predict(X_test)
y_pred_train = lasso.predict(X_train)

r2 = r2_score(y_test, y_pred)
my_dict['Lasso']=r2

# Plotting of train set
plt.title("CO2 Emissions (Training Set)")
plt.xlabel("Actual CO2 Emissions (g/km)")
plt.ylabel("Predicted CO2 Emissions (g/km)")
plt.scatter(y_train, y_pred_train, color="lightcoral", label="Training Data")
plt.plot(y_train, y_train, color="firebrick", label="Perfect Prediction")
plt.show()

# Plotting of test set
plt.title("CO2 Emissions (Testing Set)")
plt.xlabel("Actual CO2 Emissions (g/km)")
plt.ylabel("Predicted CO2 Emissions (g/km)")
plt.scatter(y_test, y_pred, color="lightcoral", label="Testing Data")
plt.plot(y_test, y_test, color="firebrick", label="Perfect Prediction")
plt.show()

# RANDOM FOREST

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
import matplotlib.pyplot as plt

X = data.drop('CO2 Emissions(g/km)', axis=1)

```

```

y = data['CO2 Emissions (g/km)']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

rf_regressor = RandomForestRegressor(n_estimators=100, random_state=42)
rf_regressor.fit(X_train, y_train)

y_pred = rf_regressor.predict(X_test)
y_pred_train = rf_regressor.predict(X_train)

r2 = r2_score(y_test, y_pred)
my_dict['Random']=r2

# Plotting of train set
plt.title("CO2 Emissions (Training Set)")
plt.xlabel("Actual CO2 Emissions (g/km)")
plt.ylabel("Predicted CO2 Emissions (g/km)")
plt.scatter(y_train, y_pred_train, color="lightcoral", label="Training Data")
plt.plot(y_train, y_train, color="firebrick", label="Perfect Prediction")
plt.show()

# Plotting of test set
plt.title("CO2 Emissions (Testing Set)")
plt.xlabel("Actual CO2 Emissions (g/km)")
plt.ylabel("Predicted CO2 Emissions (g/km)")
plt.scatter(y_test, y_pred, color="lightcoral", label="Testing Data")
plt.plot(y_test, y_test, color="firebrick", label="Perfect Prediction")
plt.show()

#Decision tree
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import r2_score, mean_squared_error
model = DecisionTreeRegressor(random_state = 42)

model.fit(X_train, y_train)
dtr_pred = model.predict(X_test)
print(dtr_pred)
print(np.sqrt(mean_squared_error(y_test,dtr_pred )))
```

```

r2=r2_score(y_test, dtr_pred)
my_dict['DTR']=r2
frames = [dtr_pred, y_test.values]
result_pred = pd.DataFrame(data=frames)
result_pred = result_pred.T
result_pred.head()

#Ridge regression

from sklearn.linear_model import Ridge
model=Ridge(alpha=0.3)
model.fit(X_train,y_train)
print('Coefficient:',model.intercept_)
print('Coeff',model.coef_)
y_ridge = model.predict(X_test)
print(np.sqrt(mean_squared_error(y_ridge,y_test)))
r2=r2_score(y_test,y_ridge)
my_dict['Ridge']=r2
frames = [y_ridge, y_test.values]
result_pred = pd.DataFrame(data=frames)
result_pred = result_pred.T
print(result_pred)
ridge_pred = result_pred.rename(columns={0: 'pred_values',
1:'real_values'})
ridge_pred['pred_values'] = (ridge_pred['pred_values'].map(lambda x:
round(x,2)))

ridge_pred['diff'] = abs(ridge_pred['real_values']
-ridge_pred['pred_values'])

print('mean diff: ', abs(ridge_pred['diff']).mean())
print(ridge_pred.head(10))

# MULTIPLE REGRESSION

import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

```

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
import matplotlib.pyplot as plt

scaler = StandardScaler()

X = data.drop('CO2 Emissions(g/km)', axis=1)
y = data['CO2 Emissions(g/km)']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

regressor = LinearRegression()

regressor.fit(X_train, y_train)
my_dict={}
y_pred = regressor.predict(X_test)
y_pred_train = regressor.predict(X_train)

r2 = r2_score(y_test, y_pred)

# LASSO REGRESSION

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Lasso
from sklearn.metrics import r2_score
import matplotlib.pyplot as plt


X = data.drop('CO2 Emissions(g/km)', axis=1)
y = data['CO2 Emissions(g/km)']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
lasso = Lasso(alpha=0.1)
lasso.fit(X_train, y_train)

y_pred = lasso.predict(X_test)
y_pred_train = lasso.predict(X_train)

r2 = r2_score(y_test, y_pred)
my_dict['Lasso']=r2

# RANDOM FOREST

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score
import matplotlib.pyplot as plt

X = data.drop('CO2 Emissions(g/km)', axis=1)
y = data['CO2 Emissions(g/km)']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

rf_regressor = RandomForestRegressor(n_estimators=100, random_state=42)
rf_regressor.fit(X_train, y_train)

y_pred = rf_regressor.predict(X_test)
y_pred_train = rf_regressor.predict(X_train)

r2 = r2_score(y_test, y_pred)
my_dict['Random']=r2

#Decision tree
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import r2_score,mean_squared_error
model = DecisionTreeRegressor(random_state = 42)

model.fit(X_train, y_train)
dtr_pred = model.predict(X_test)
```

```

print(dtr_pred)
print(np.sqrt(mean_squared_error(y_test,dtr_pred)))
r2=r2_score(y_test, dtr_pred)
my_dict['DTR']=r2
frames = [dtr_pred, y_test.values]
result_pred = pd.DataFrame(data=frames)
result_pred = result_pred.T
result_pred.head()

#Ridge regression

from sklearn.linear_model import Ridge
model=Ridge(alpha=0.3)
model.fit(X_train,y_train)
print('Coefficient:',model.intercept_)
print('Coeff',model.coef_)
y_ridge = model.predict(X_test)
print(np.sqrt(mean_squared_error(y_ridge,y_test)))
r2=r2_score(y_test,y_ridge)
my_dict['Ridge']=r2
frames = [y_ridge, y_test.values]
result_pred = pd.DataFrame(data=frames)
result_pred = result_pred.T
print(result_pred)
ridge_pred = result_pred.rename(columns={0: 'pred_values',
1:'real_values'})
ridge_pred['pred_values'] = (ridge_pred['pred_values'].map(lambda x:
round(x,2)))

ridge_pred['diff'] = abs(ridge_pred['real_values']
-ridge_pred['pred_values'])

print('mean diff: ', abs(ridge_pred['diff']).mean())
print(ridge_pred.head(10))

```

OUTPUT:

```

Make          0
Model         0
Vehicle Class 0
Engine Size(L) 0
Cylinders     0
Transmission   0
Fuel Type      0
Fuel Consumption City (L/100 km) 0
Fuel Consumption Hwy (L/100 km) 0
Fuel Consumption Comb (L/100 km) 0
Fuel Consumption Comb (mpg)      0
CO2 Emissions(g/km)            0
dtype: int64

```

There are 1103 duplicated rows in the data

	count	mean	std	min	25%	50%	75%	max
Engine Size(L)	6282.0	3.161812	1.365201	0.9	2.0	3.0	3.7	8.4
Cylinders	6282.0	5.618911	1.846250	3.0	4.0	6.0	6.0	16.0
Fuel Consumption City (L/100 km)	6282.0	12.610220	3.553066	4.2	10.1	12.1	14.7	30.6
Fuel Consumption Hwy (L/100 km)	6282.0	9.070583	2.278884	4.0	7.5	8.7	10.3	20.6
Fuel Consumption Comb (L/100 km)	6282.0	11.017876	2.946876	4.1	8.9	10.6	12.7	26.1
Fuel Consumption Comb (mpg)	6282.0	27.411016	7.245318	11.0	22.0	27.0	32.0	69.0
CO2 Emissions(g/km)	6282.0	251.157752	59.290426	96.0	208.0	246.0	289.0	522.0

Make	object
Model	object
Vehicle Class	object
Engine Size(L)	float64
Cylinders	int64
Transmission	object
Fuel Type	object

```

Fuel Consumption City (L/100 km)      float64
Fuel Consumption Hwy (L/100 km)        float64
Fuel Consumption Comb (L/100 km)       float64
Fuel Consumption Comb (mpg)           int64
CO2 Emissions (g/km)                  int64
dtype: object

Empty DataFrame
Columns: [Make, Model, Vehicle Class, Engine Size(L), Cylinders,
Transmission, Fuel Type, Fuel Consumption City (L/100 km), Fuel
Consumption Hwy (L/100 km), Fuel Consumption Comb (L/100 km), Fuel
Consumption Comb (mpg), CO2 Emissions (g/km)]
Index: []

Make                           True
Model                          True
Vehicle Class                  True
Engine Size(L)                 False
Cylinders                      False
Transmission                   True
Fuel Type                      True
Fuel Consumption City (L/100 km) False
Fuel Consumption Hwy (L/100 km)  False
Fuel Consumption Comb (L/100 km) False
Fuel Consumption Comb (mpg)     False
CO2 Emissions (g/km)           False
dtype: bool

Fuel Consumption Comb (mpg)      -0.906783
Cylinders                       0.834687
Engine Size(L)                  0.854802
Fuel Consumption Hwy (L/100 km)  0.883424
Fuel Consumption Comb (L/100 km) 0.916840
Fuel Consumption City (L/100 km) 0.918756
CO2 Emissions (g/km)             1.000000
Name: CO2 Emissions (g/km), dtype: float64

#checking for distinct model names
array(['ILX', 'ILX HYBRID', 'MDX 4WD', ...,
       'Tacoma 4WD D-Cab TRD Off-Road/Pro', 'Atlas Cross Sport 4MOTION',
       'XC40 T4 AWD'], dtype=object)

array(['ACURA', 'ALFA ROMEO', 'ASTON MARTIN', 'AUDI', 'BENTLEY', 'BMW',
       'BUICK', 'CADILLAC', 'CHEVROLET', 'CHRYSLER', 'DODGE', 'FIAT',
       'FORD', 'GMC', 'HONDA', 'HYUNDAI', 'INFINITI', 'JAGUAR', 'JEEP',
       'KIA', 'LAMBORGHINI', 'LAND ROVER', 'LEXUS', 'LINCOLN', 'MASERATI'],
      dtype=object)

```

```
'MAZDA', 'MERCEDES-BENZ', 'MINI', 'MITSUBISHI', 'NISSAN',
'PORSCHE', 'RAM', 'ROLLS-ROYCE', 'SCION', 'SMART', 'SRT', 'SUBARU',
'TOYOTA', 'VOLKSWAGEN', 'VOLVO', 'GENESIS', 'BUGATTI'],
dtype=object)

['Automatic Selection' 'Manual' 'Autonomous' 'Automated Manual'
 'Automatic']
```

Transmission

```
Automatic Selection      3127
Automatic                 1851
Manual                     1185
Automated Manual            646
Autonomous                  576
Name: count, dtype: int64
```

```
['Premium gasoline' 'Diesel' 'Regular gasoline' 'Ethanol' 'Natural gas']
['COMPACT' 'SUV - SMALL' 'MID-SIZE' 'TWO-SEATER' 'MINICOMPACT'
 'SUBCOMPACT' 'FULL-SIZE' 'STATION WAGON - SMALL' 'SUV - STANDARD'
 'VAN - CARGO' 'VAN - PASSENGER' 'PICKUP TRUCK - STANDARD' 'MINIVAN'
 'SPECIAL PURPOSE VEHICLE' 'STATION WAGON - MID-SIZE'
 'PICKUP TRUCK - SMALL']
```

Fuel Type

```
Regular gasoline   3039
Premium gasoline   2765
Ethanol             330
Diesel              147
Natural gas          1
Name: count, dtype: int64
```

	Make	Model	Vehicle Class	Engine Size(L)	Cylinders
\					
2439	CHEVROLET	IMPALA DUAL FUEL	MID-SIZE	3.6	6

	Transmission	Fuel Type	Fuel Consumption City (L/100 km)	\
2439	AS6	Natural gas	15.2	

	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	\
2439	9.5	12.7	

	Fuel Consumption Comb (mpg)	CO2 Emissions(g/km)
2439	22	213

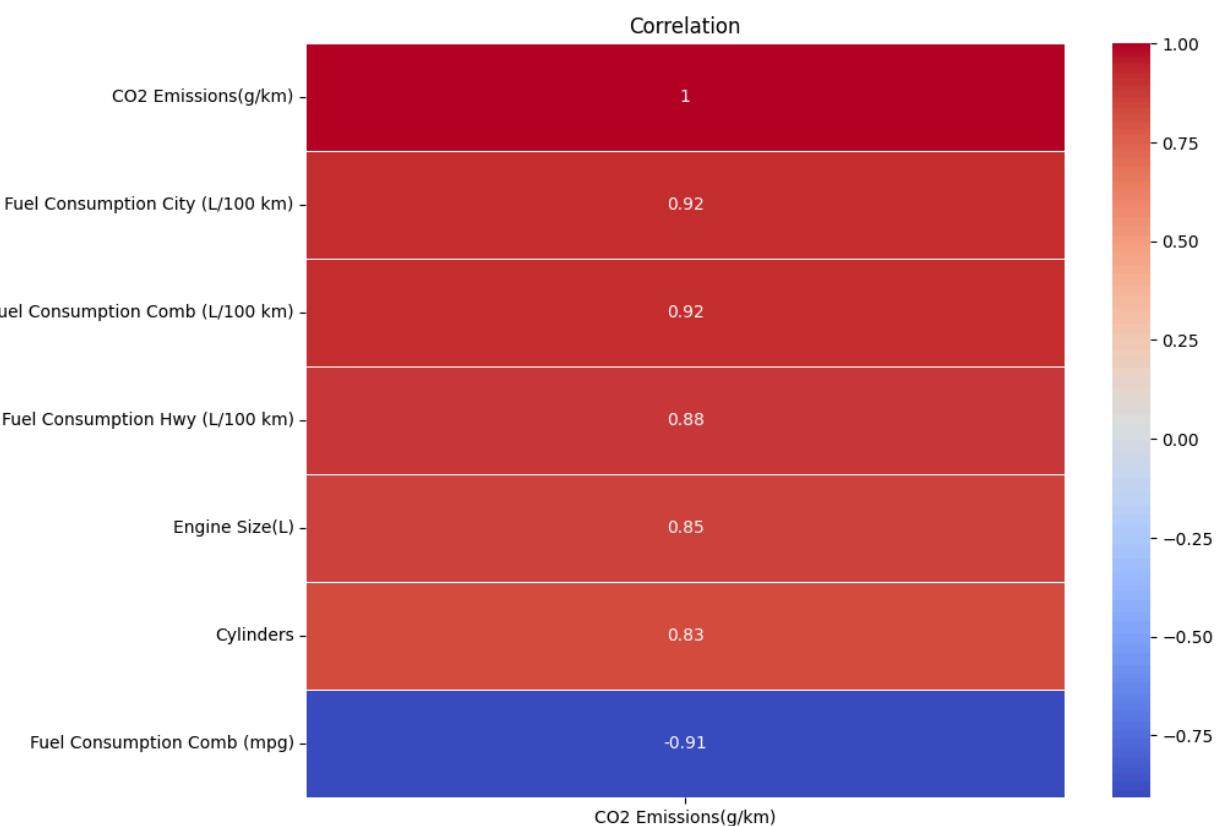
	Fuel_Ethanol	Fuel_Premium gasoline	Fuel-Regular gasoline
0	False	True	False
1	False	True	False
2	False	True	False
3	False	True	False
4	False	True	False

5 rows × 41 columns

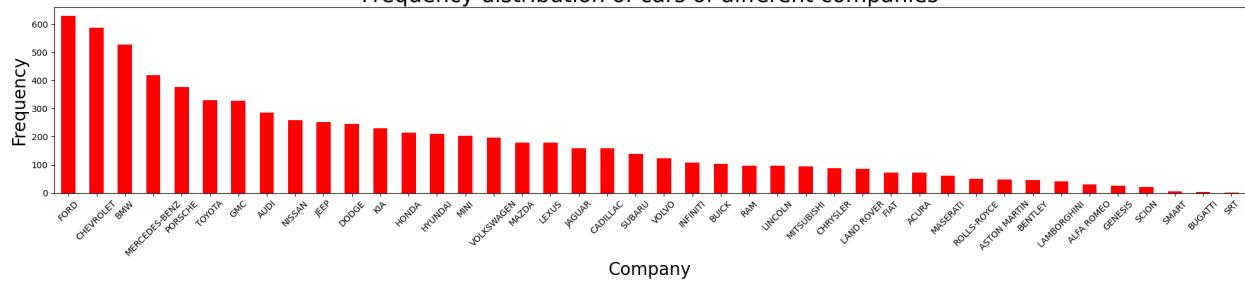
	5	1	1	3.5	6	12.7	9.1	11.1	25	F	.	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F				
1		0								a	.	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a		
3		0								l	.	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l			
		6								s		s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s		
										e		e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	
	5	4	1	3.5	6	12.1	8.7	10.6	27	F	.	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F				
1		0								a	.	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a	a		
4		0								l	.	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l	l			
		6								s		s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s	s		
										e		e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	e	

5 rows × 38 columns

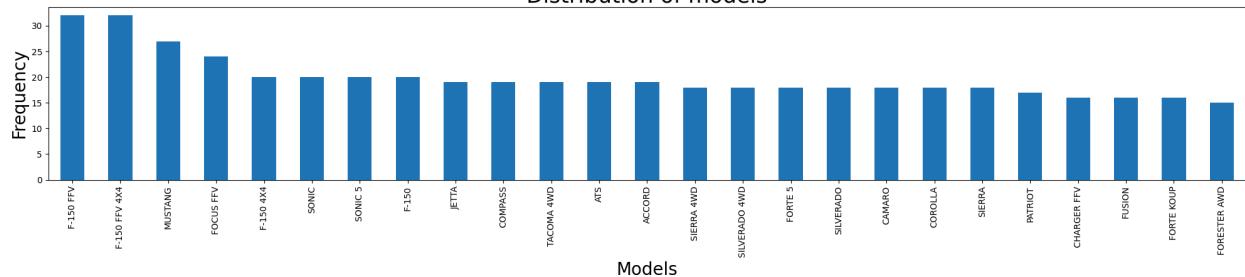
DATA VISUALIZATION:



Frequency distribution of cars of different companies



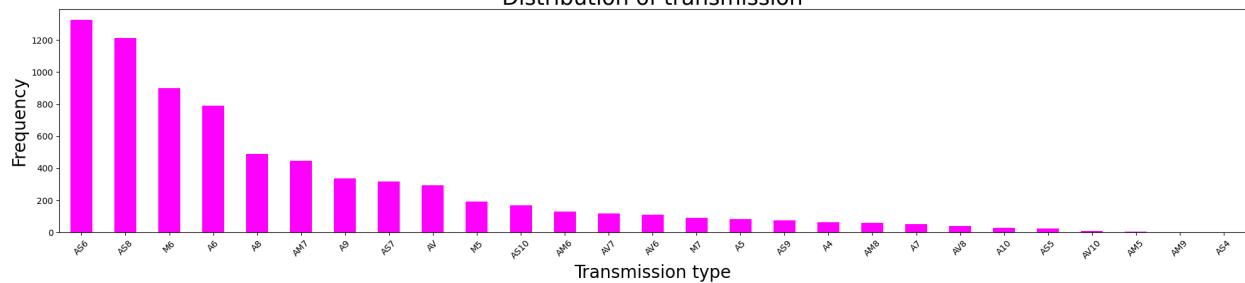
Distribution of models



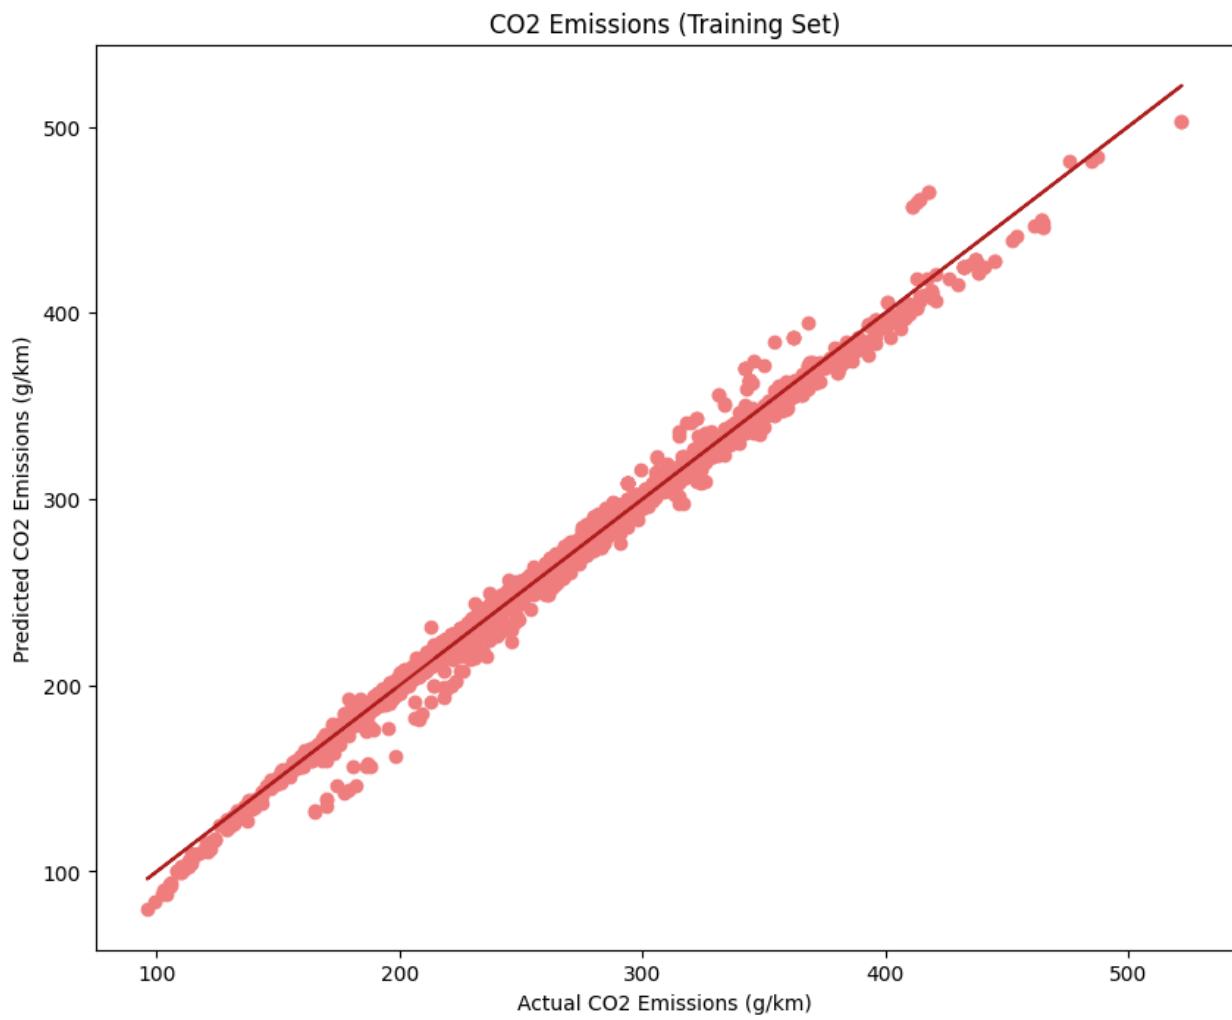
Vehicle class distribution

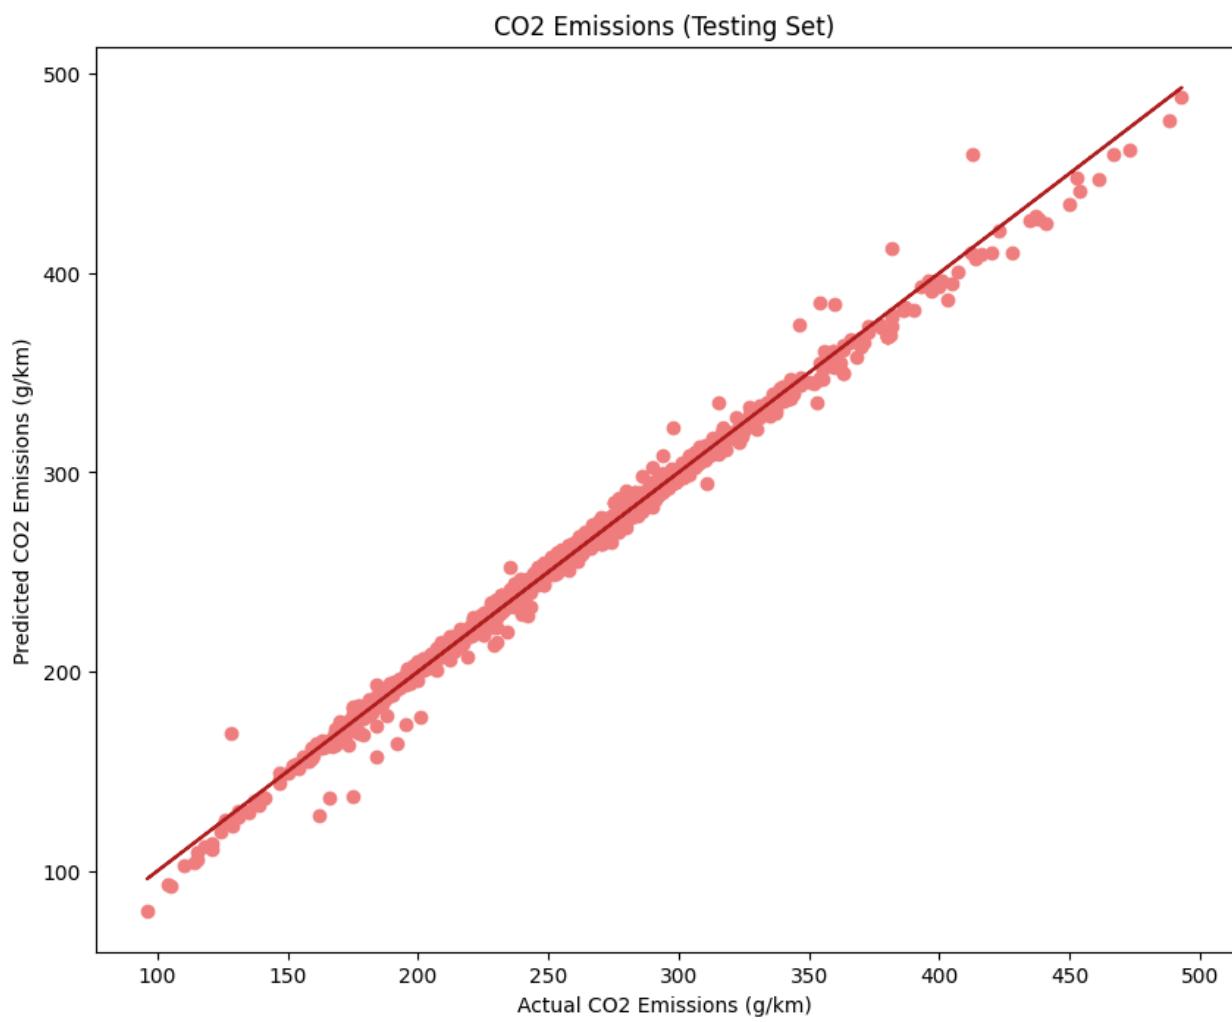


Distribution of transmission



MULTIPLE REGRESSION





127.0.0.1:5000

Vehicle Specification

Company Name: Mahindra

Model Name: XUV

Vehicle Class: Car

Engine Size (L): 2.18

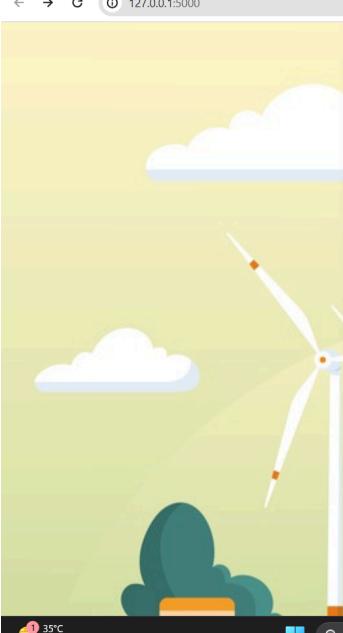
Cylinders: 3

Transmission: Manual

Fuel Type: Diesel

Fuel Consumption in City (L/100 km): 7

Fuel Consumption on Highway (L/100 km): 5



127.0.0.1:5000

Engine Size (L): 2.18

Cylinders: 3

Transmission: Manual

Fuel Type: Diesel

Fuel Consumption in City (L/100 km): 7

Fuel Consumption on Highway (L/100 km): 5

Average Fuel Consumption On Highway & City (L/100 km): 6

Fuel Consumption Rate (mpg): 20

Submit



127.0.0.1:5000/newpage

Vehicle CO2 Emission Evaluation

The vehicle is considered **BAD** in terms of CO2 emission.

Machine Learning Algorithms for Prediction

Algorithm	Predicted CO2 Emission	Accuracy
Multiple Regression	225.9	90.18902750130587
Lasso Regression	225.89	90.19024914845113
Ridge Regression	225.88	90.18916736931757
Decision Tree	150.33	95.94480830409137
Random Forest	176.24	97.48088731502416

35°C Haze

ENG IN 16:56 17-04-2024

127.0.0.1:5000/newpage

Vehicle CO2 Emission Evaluation

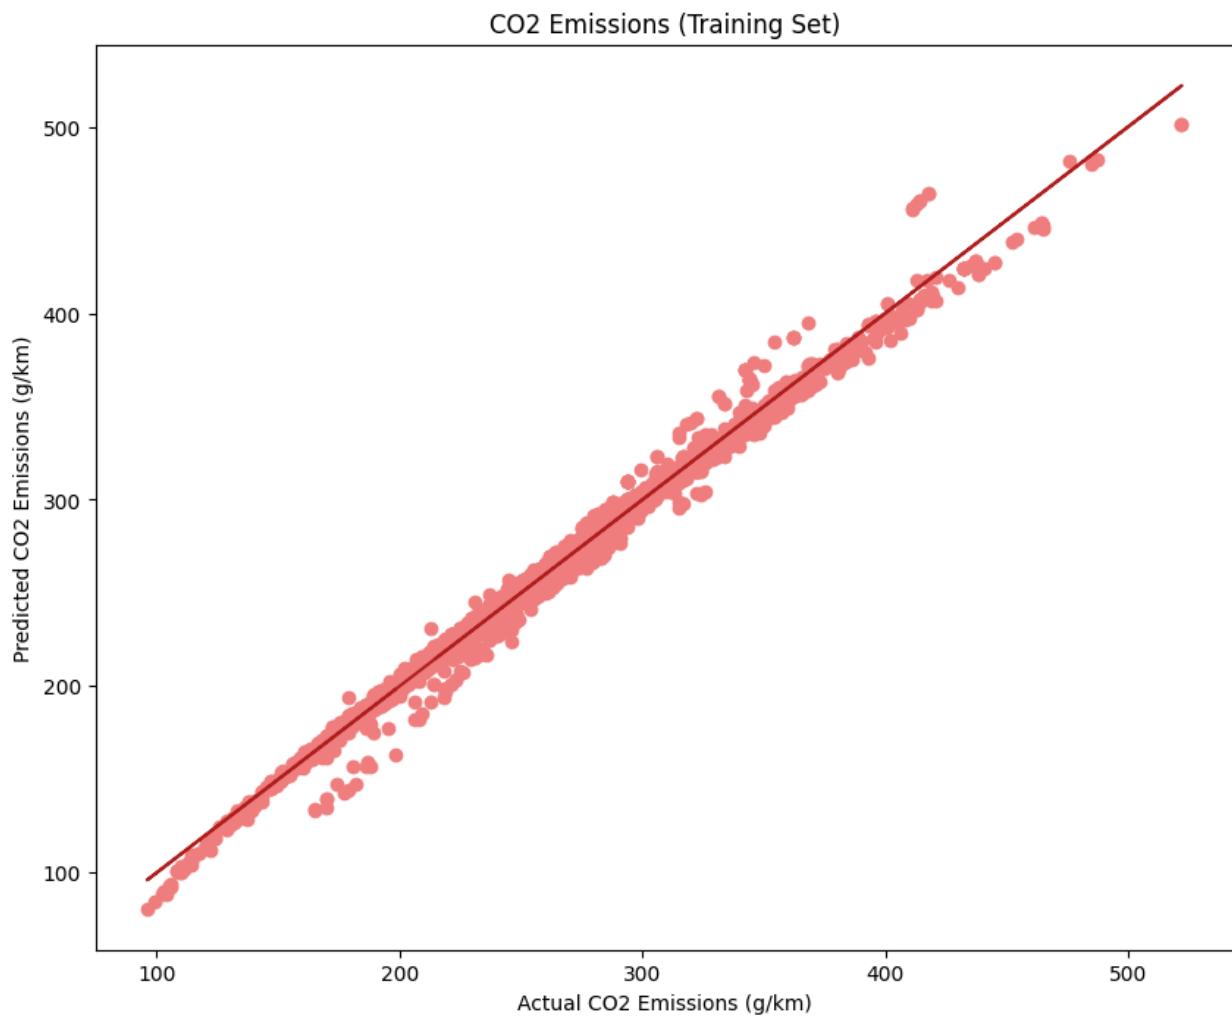
The vehicle is considered **GOOD** in terms of CO2 emission.

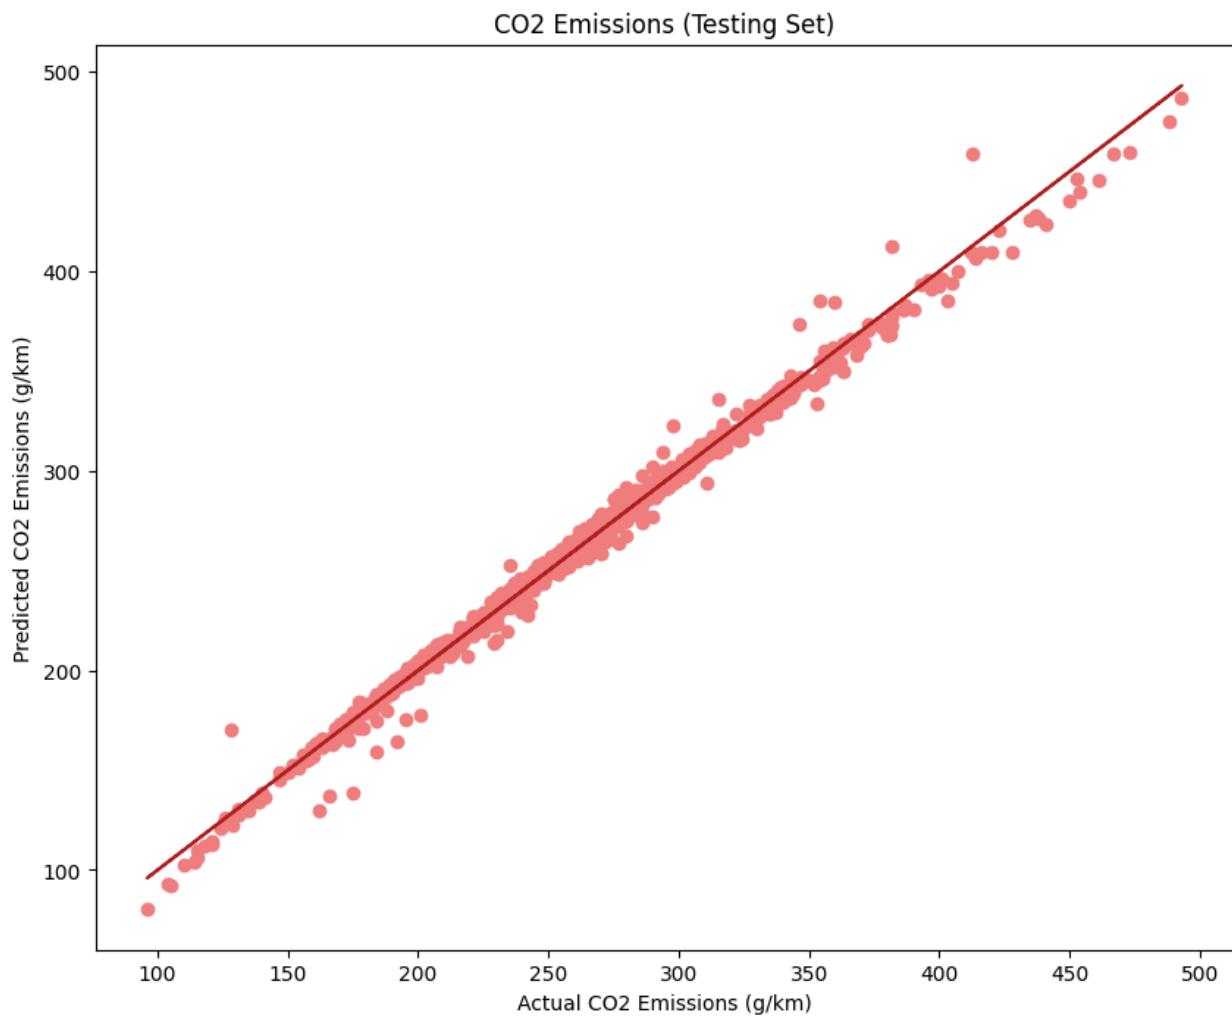
Machine Learning Algorithms for Prediction

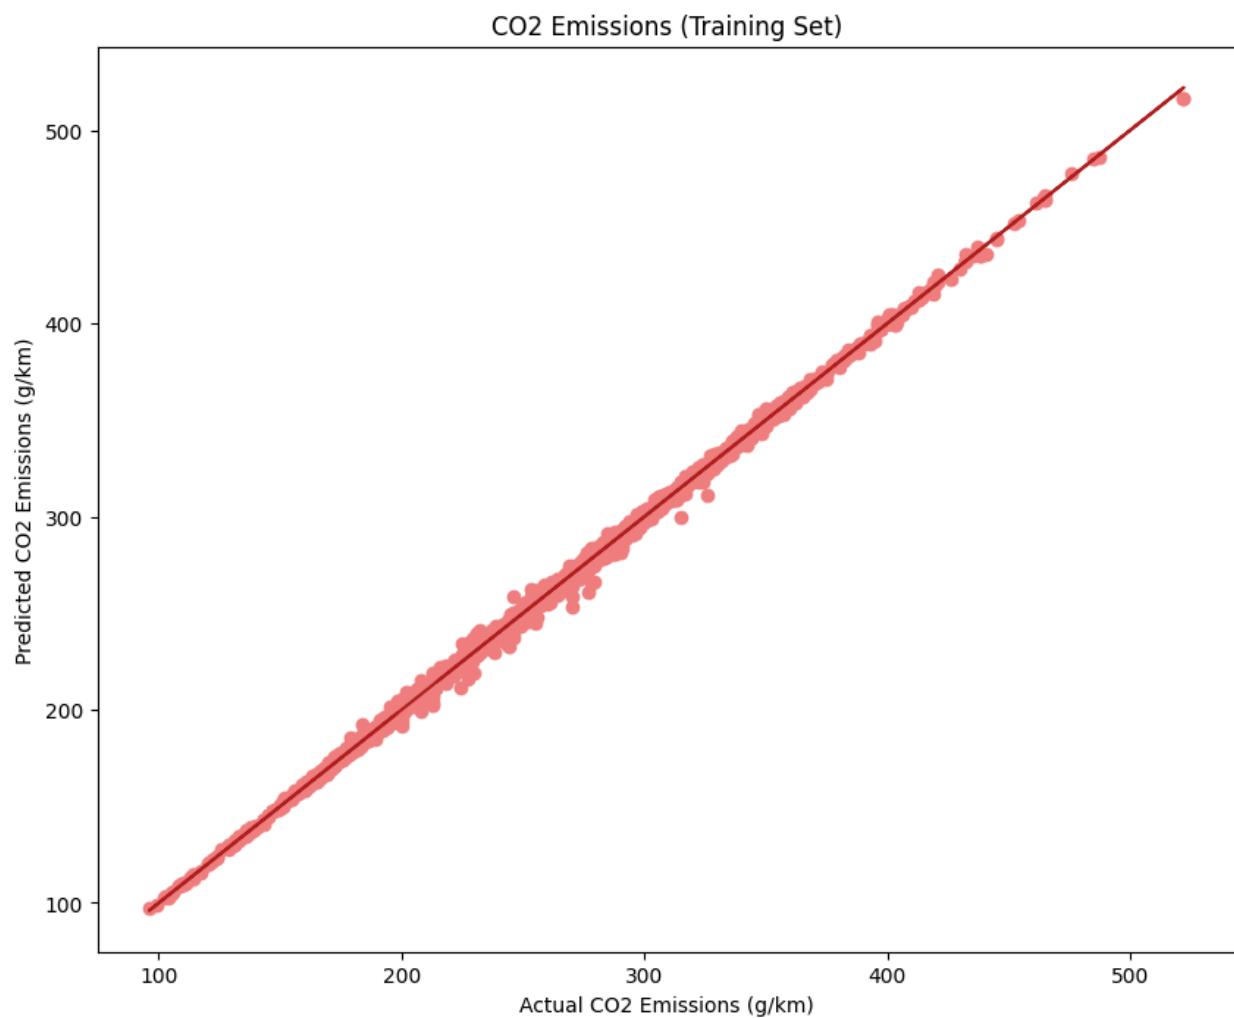
Algorithm	Predicted CO2 Emission	Accuracy
Multiple Regression	160.37	90.18902750130587
Lasso Regression	160.56	90.19024914845113
Ridge Regression	160.36	90.18916736931757
Decision Tree	145.5	96.19924682225364
Random Forest	149.4	97.63699954154117

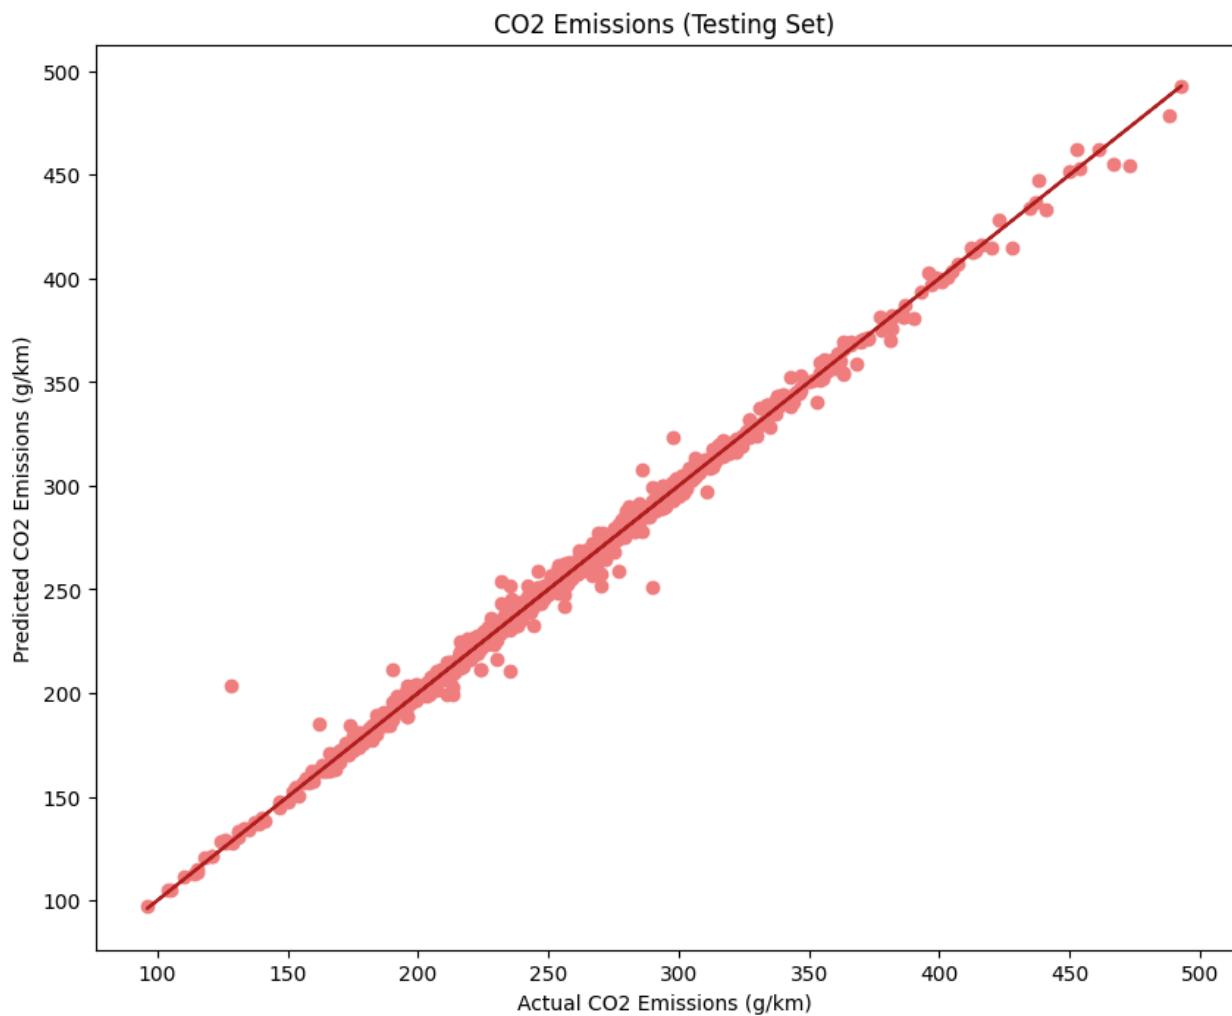
35°C Haze

ENG IN 16:55 17-04-2024









CONCLUSION:

Provide recommendations for policymakers, manufacturers, and consumers based on the model results, suggesting strategies for reducing emissions and promoting environmental sustainability in the transportation sector.